

Lecture 14: Device Management in Operating Systems

Operating Systems 2025

Prof. A.I. Companion

June 27, 2025

Hello everyone, and welcome. This document is a complete rewrite of our lecture on Device Management. It is designed for any student who missed the class and needs to understand all the material as if they were present. We will cover every concept, explanation, and example in the order it was presented.

Introduction and Administrative Notes

As is our custom at the beginning of class, we started with a quick name check to take attendance. After that, we moved on to some important announcements regarding our upcoming exam before diving into today's main topic.

Important: Exam 1 Information

Before we begin with new material, let's go over the details for our first exam. Please pay close attention to this, as it contains critical information for your preparation.

The exam is scheduled for **Friday, April 4th, 2025**, from **9:00 AM to 12:00 PM**. The location will be room **3B2308**.

What to Bring:

- * **2B Pencils:** You absolutely must bring 2B pencils for the multiple-choice section of the exam. The answer sheets are graded by a machine, and if you use a different type of pencil or a pen, the machine cannot read your answers. This could result in a score of zero for that section, so please do not forget this. It's a good idea to bring one or two, along with an eraser.
- * **Pens and Correction Fluid:** For the written-answer sections, you can use pens. You may also bring correction fluid (like liquid paper) in case you need to make changes to your answers.
- * **Paper Dictionary:** You are permitted to use a paper dictionary. This is your choice; you can bring one or not, but electronic dictionaries are not allowed.
- * **One A4 Paper Note (Cheat Sheet):** The exam is closed-book, but you are allowed to bring **one single A4-sized sheet of paper** with notes. You can prepare this note however you like—it can be handwritten, or you can type it up on a computer using software like OneNote and print it out. You can use both sides of the paper.
- * **Crucial Rule:** You **must** write your full name and student ID on this note sheet. When you enter the exam room, this sheet will be considered part of your exam submission. You can use it during the exam, but you must turn it in with your exam paper at the end. You cannot take it out of the room with you. Is that clear? Please do not forget to write your name and ID on it.

Exam Question Format:

Be aware that some questions, particularly in the multiple-choice section, may have **more than one correct answer**. These questions will be clearly indicated with phrases like "choose all that apply" or "select all that apply." When you see this, it's a signal that you should look for multiple valid options.

Exam Scope and Syllabus Clarification

Let's clarify the exact topics that will be covered on the exam. The scope includes the material from Week 2 through Week 10, but with a slight adjustment from the original syllabus.

Here is the breakdown of the topics covered:

1. **Week 2:** Introduction to Operating Systems
2. **Week 3:** Early Memory Management
3. **Weeks 4, 5, 6:** Advanced Memory Management (Note: Week 6 was an exam break)
4. **Weeks 7, 8, 9, 10:** Process Management and Concurrency

You might notice a delay compared to the original course syllabus. We had initially planned to finish the topic of Concurrency by Week 8. However, we needed more time to cover it thoroughly, so we spent Weeks 9 and 10 on Concurrency as well. This means there was a two-week delay in our schedule.

Therefore, the exam scope will cover these three main pillars:

1. **Introduction to Operating Systems**
2. **Memory Manager**
3. **Process Manager (including Concurrency)**

Today's topic, **Device Manager**, will **not** be on this exam. Hopefully, everyone is now clear on the exam scope. If you have any confusion, please let me know.

Lecture: Device Management

Alright, let's move on to the third major component of the operating system we'll be studying: the **Device Manager**. So far, we have discussed two other key managers: the **Memory Manager**, which handles the allocation and deallocation of memory, and the **Process Manager**, which is responsible for creating, scheduling, and terminating processes. The Device Manager is the third piece of this puzzle.

Learning Objectives

By the end of this lecture (and the next), you should be able to:

- * Compare and contrast dedicated, shared, and virtual devices.
- * Understand how blocking and buffering can improve I/O (Input/Output) performance.
- * Calculate **seek time**, **search time**, and **transfer time** for a hard disk drive. This is a practical skill; if you were to design a new I/O device, you'd need to be able to calculate its performance metrics.
- * Explain how access times differ across various types of storage devices.
- * Describe the strengths and weaknesses of common disk seek strategies.
- * Understand the different levels of RAID (Redundant Array of Independent Disks). This topic is marked as optional for now; we may or may not get to it depending on our timing. The plan is to finish the core concepts of the Device Manager this week and potentially continue with RAID next week if time permits. Next, we will move on to the Network Manager, another of the five key managers in an OS.

Today's Agenda

Here is the plan for today's discussion:

1. **Device Management Overview:** What it is and why we need it.
 2. **Access Time:** A fundamental concept for measuring device performance.
 3. **Sequential Access Storage Media:** We'll explore whether technologies like magnetic tape are still relevant today, especially in the context of cloud computing.
 4. **Magnetic Disk Storage:** The inner workings of hard drives.
 5. **Device Handler Seek Strategies:** We'll look at algorithms used to optimize disk performance. This will be a hands-on part of the lecture with examples.
 6. **Optical Disk Storage, Flash Memory, and Solid-State Devices (SSDs).**
 7. **The I/O Subsystem.**
 8. **RAID:** As mentioned, we will likely cover this in our next session.
-

Part 1: The Role of the Device Manager

So, what exactly is the **Device Manager**? In simple terms, your operating system needs a way to communicate with and control all the hardware components connected to your computer. These components are called **peripheral devices**. This is precisely the job of the Device Manager.

The Device Manager is the part of the operating system that controls all input and output (I/O) devices. This includes everything from your keyboard, mouse, and printer to joysticks and USB drives. It uses specialized software components called **device controllers** to manage these devices, checking their status (e.g., is the printer ready?) and handling access requests from different applications.

A key function is to allow application software (like your web browser or a game) to communicate with these hardware devices without needing to know the specific, complex details of how each device works.

The Four Basic Functions of the Device Manager

The Device Manager has four primary responsibilities:

1. **Monitoring Device Status:** It constantly keeps track of the status of every device. Is the hard drive busy? Is the monitor on? Is a USB drive connected? This applies to storage hardware like hard disk drives (HDDs) and solid-state drives (SSDs), as well as monitors, scanners, and other peripherals.
2. **Enforcing Policies:** The Device Manager implements policies to determine how devices are allocated. These policies decide which process gets to use a device, when it gets it, and for how long. This is very similar to the scheduling policies we saw in the Process Manager, such as First-Come, First-Served (FCFS) or Shortest Job Next. You can apply similar logic to decide how to serve I/O requests.
3. **Allocating Devices:** When a process requests a device, the manager allocates it appropriately.
4. **Deallocating Devices:** Once a process is finished with a device, the manager must deallocate it, or "take it back," so it becomes available for other processes. Deallocation can happen at two levels:
 - **Process/Task Level:** A device is temporarily released after an I/O command has been executed.
 - **Job Level:** A device is permanently released when the entire job (which may consist of multiple processes) is finished.

If you use a Windows machine, you can see this in action by opening the "Device Manager" utility. It provides a list of all the hardware connected to your system and is the user-facing part of the OS component we are discussing.

Part 2: Types of Devices

The operating system manages both physical hardware and **virtual devices**. A virtual device is a piece of software that mimics a physical device, like a virtual printer that "prints" to a PDF file instead of paper.

We can categorize devices into three main types based on how they are shared:

1. **Dedicated Devices**
2. **Shared Devices**
3. **Virtual Devices**

1. Dedicated Devices

A **dedicated device** is a device that is assigned to a single process for the entire duration of its task. It is not released until that process is completely finished.

- **Why it matters:** This type of allocation is necessary for devices where it would be inefficient or nonsensical to have multiple processes using them at once.
- **Examples:** Printers, tape drives, monitors, and mice are classic examples. When you print a document, you want your entire document to print without another user's pages getting mixed in.
- **Analogy:** Imagine sharing a single monitor between two users simultaneously. One half of the screen shows your video game, while the other half shows someone else's Netflix movie. It would be an awkward and unusable experience. This is why a monitor is dedicated to one user (or one primary process) at a time. A dedicated device serves one job until that job is completely done. Even if the job isn't using the device 100% of the time, the device remains allocated and cannot be given to another process.

2. Shared Devices

A **shared device** can be used by multiple processes or users concurrently.

- **Why it matters:** Sharing allows for more efficient use of resources, as multiple processes can access the device without having to wait for one another to finish completely.
- **Example:** A **Direct Access Storage Device (DASD)**, like a hard disk drive, is a perfect example. Because a hard drive can randomly access any location on its platters, it can handle requests from multiple processes by interleaving them. The Device Manager controls this access, resolving conflicts and scheduling requests based on its preset policies.

3. Virtual Devices

A **virtual device** is a clever combination of dedicated and shared concepts. It's a way to make a dedicated device *behave* like a shared device, improving efficiency.

- **Why it matters:** Virtual devices can overcome the limitations of slow, dedicated hardware, making the system feel faster and more responsive to the user.
- **The Spooling Analogy:** The most common example is a printer **spooler**. A physical printer is a dedicated device—it can only print one document at a time. If multiple users send print jobs, they would have to wait in a queue. **Spooling** (an acronym for Simultaneous Peripheral Operations On-Line) solves this.
 - When you click “print,” your document isn't sent directly to the physical printer. Instead, it's sent to a special area on a fast secondary storage device (like a hard drive). This is the “spool.”
 - The operating system manages this spool, and your application is immediately freed up. You can go back to work without waiting for the physical printing to finish.
 - The spooling program then feeds the jobs from the disk to the printer one by one, whenever the printer is ready.
 - In this way, the slow, dedicated printer is converted into a fast, sharable virtual device. The hard drive acts as a buffer, making the whole process much more efficient.

A Closer Look at USB Devices

Everyone uses USB (Universal Serial Bus). A single USB port, often expanded with a hub, can support many devices simultaneously—flash drives, phones, keyboards, etc.

Communication happens through logical “pipes” that link the main computer (the **host controller**) to the endpoint devices. When you plug in a USB device, the host controller assigns it a unique ID and manages data exchange.

Interestingly, the host controller assigns **bandwidth priority** to different types of data:

- * **High Priority:** Given to real-time, uninterrupted data streams like audio and video. If you're using a USB microphone or webcam, it gets top priority to ensure smooth performance.
- * **Medium Priority:** Assigned to devices that send data occasionally, like a keyboard. The host controller will interrupt other tasks if needed to process your keystrokes.
- * **Low Priority:** Used for bulk data transfers, like copying files to or from a flash drive.

This priority system ensures that time-sensitive tasks (like a video call) are not disrupted by less urgent tasks (like a background file transfer). If you're copying a large file (low priority) and start typing on your keyboard (medium priority), the system will prioritize processing your keystrokes.

Part 3: Managing I/O Requests

An **I/O device** is any piece of hardware that allows a computer to communicate with the outside world. Inputs are data received (e.g.,

from a keyboard), and outputs are data sent (e.g., to a monitor). Some devices, like modems, network cards, or touch screens, handle both input and output.

For the operating system to talk to a device, it needs a **device driver**. A device driver is a specialized piece of software that acts as a translator. It converts the generic commands from the OS into specific instructions that a particular piece of hardware understands. In modern operating systems, many common device drivers come pre-installed.

The Device Manager breaks down the task of managing an I/O request into three parts, handled by different components:

1. **The I/O Channel:** A specialized, programmable unit that sits between the CPU and the device's control unit. Its job is to synchronize the fast speed of the CPU with the much slower speed of the I/O device.
2. **The Control Unit:** This is hardware, often built into the device itself, that contains the electronic components needed to operate a specific type of device (e.g., a disk controller for all hard drives).
3. **The Device Itself.**

To coordinate all of this, the operating system uses several key software components:

- **I/O Traffic Controller:** This component watches the status of all I/O devices, control units, and channels. Its main tasks are:
 1. To determine if a path to a device is available.
 2. If multiple paths are available, to select the best one.
 3. If all paths are busy, to determine when one will become free.Think of it like a road traffic controller managing intersections and directing cars (data) along the best routes.
- **I/O Scheduler:** Just like the process scheduler, the I/O scheduler implements policies to decide which I/O request gets access to a device, control unit, or channel, and in what order. Most I/O tasks are **non-preemptive**, meaning they run to completion once started. However, some advanced systems allow for preemption by breaking I/O tasks into smaller stages to handle urgent requests. Your design choices as a system architect determine whether the system is preemptive or not.
- **I/O Device Handler:** This component performs the actual transfer of data and processes device interrupts (signals from the device that it has completed a task or needs attention). Each type of device has its own unique handler.

The Control Block System

To keep track of everything, the I/O Traffic Controller uses a database of “control blocks,” which are data structures that store the status and information for each component. You might remember Process Control Blocks (PCBs) and Thread Control Blocks (TCBs) from our process management lectures; this is the same idea applied to I/O.

1. Channel Control Block (CCB):

- **What it is:** Represents an I/O channel, which is the data lane between the CPU/main memory and the I/O devices.
- **What it stores:** Channel ID, status (busy/free), a list of control units connected to it, and a list of processes waiting for the channel.
- **Analogy:** If data packets are cars, the channel is the road lane. The Channel Control Block is like the traffic light system that controls the flow of cars in that lane.

2. Control Unit Control Block (CUCB):

- **What it is:** Represents an I/O control unit, which manages a specific *type* of device (e.g., one controller for all printers, another for all disk drives).
- **What it stores:** Control unit ID, status (active/waiting/available), the channel it's connected to, a list of devices it controls, and processes waiting for it.
- **Analogy:** An I/O control unit is like a construction site manager who oversees a group of similar workers (the devices). The CUCB is the manager's clipboard, keeping track of all the workers and their tasks.

3. Device Control Box (DCB):

- **What it is:** Represents a single, individual I/O device.
- **What it stores:** Device ID, status, location, the control unit it's connected to, and a list of jobs currently using it.
- **Analogy:** A DCB is like a profile card for each individual worker (device). It shows their current status and what specific task they are working on.

This creates a clear hierarchy of communication: a request from the **CPU** goes through a **Channel** to a **Control Unit**, which then directs the command to a specific **Device**.

I/O in the Cloud Computing Era

Today, with cloud computing, the nature of I/O is changing. You can now access an I/O device, like a specialized scientific sensor, that is physically located in another country. The role of the local operating system expands to include managing access to these remote devices over the internet. For instance, to access data from a weather station in Africa, your local OS would communicate with the operating system running on the remote machine in Africa.

Part 4: Device Access Methods

There are two fundamental ways to access data on a storage device:

1. **Sequential Access Method:** Data is accessed in a linear sequence, one piece after another. To get to record #6, you must first read through records #1, #2, #3, #4, and #5.
2. **Random (or Direct) Access Method:** You can jump directly to any data location without reading through the preceding data. Our main memory (RAM - Random Access Memory) is a prime example of this.

Sequential Access Storage Media: The Magnetic Tape

The classic example of a sequential access device is the **magnetic tape**. While it may seem like ancient technology, it was one of the earliest forms of secondary storage and, believe it or not, is still used today for specific purposes.

A magnetic tape stores data on a long strip of magnetic material. To read or write, the tape must physically move past a **read/write head**. This makes accessing data very slow if the data you need is far down the tape.

- **Structure:** Data is typically stored in nine parallel **tracks**. Eight tracks are for the data itself (representing one byte), and the ninth track is for a **parity bit**. The parity bit is used for basic error checking. If a small part of the tape gets damaged, the parity system can detect that an error has occurred in that column of bits.
- **Modern Use Case:** Because tapes are very cheap and can store massive amounts of data (thousands of bytes per inch), they are still used for mainframe backups and archival storage. Cloud services like **Amazon Web Services (AWS) Glacier** use tape storage for data that is rarely accessed and doesn't need to be retrieved quickly. It's a cost-effective solution for "cold storage."

Improving Efficiency: Inter-Record Gaps (IRG) and Blocking

When storing individual records on a tape, a physical gap is needed between them. This is called an **Inter-Record Gap (IRG)**. The tape drive needs this gap to have enough time and space to stop after reading a record before starting to move again for the next one.

- **The Problem:** These gaps waste a lot of space. For example, a single record might only take up 0.1 inches of tape, but the IRG next to it could be 0.5 inches. This means you are using more tape for empty space than for your actual data! If you have 10 records, you'll have 9 gaps, which adds up to a significant amount of wasted tape.
- **The Solution: Blocking.** To make this more efficient, we can group multiple records together into a **block**. Instead of having a gap after every single record, we only have one gap, now called an **Inter-Block Gap (IBG)**, after the entire block.
 - **Advantages of Blocking:**
 1. **Fewer I/O Operations:** A single "read" command can bring the entire block of records into memory at once.
 2. **Less Wasted Tape:** The amount of tape used for data is much larger than the amount used for gaps.
 - **Disadvantages of Blocking:**
 1. **Overhead:** Software routines are needed to handle the blocking and de-blocking of records.
 2. **Wasted Buffer Space:** The entire block must be read into a buffer in main memory, even if the application only needs one record from that block. This can waste memory.

The performance of a tape drive is measured by its **transfer rate**.

Transfer Rate = Tape Density (bytes per inch) × Transport Speed (inches per second)

For example, a tape with a density of 2,000 bytes/inch and a transport speed of 200 inches/second has a transfer rate of 400,000 bytes/second (or 400 KB/s).

Part 5: Direct Access Storage Devices (DASD)

Now let's turn to **Direct Access Storage Devices (DASD)**, which allow the system to read or write data at any specific location without processing the data in sequence. This is also known as random access.

Examples include:

- * Magnetic Disks (Hard Disk Drives - HDDs)
- * Optical Disks (CDs, DVDs, Blu-ray)
- * Flash Memory and Solid-State Drives (SSDs)

Access time on these devices can vary depending on the physical location of the data, but it is significantly faster than sequential access on a magnetic tape.

Anatomy of a Magnetic Disk Drive (HDD)

A hard disk drive consists of one or more rigid metal platters coated with a magnetic material.

- * **Platters and Surfaces:** Each platter has two recording surfaces (top and bottom).
- * **Tracks:** Each surface is divided into thousands of concentric circles called **tracks**. Tracks are numbered starting from 0 at the outermost

edge.

* **Read/Write Heads:** A **read/write head** is used to access the data on each surface. These heads are attached to an **arm assembly** that moves them across the tracks.

* **Cylinder:** A **cylinder** is a virtual concept. It is the set of all tracks that are at the same distance from the center, across all platters. For example, track 36 on every single platter surface forms Cylinder 36. The number of cylinders on a disk is equal to the number of tracks on a single surface.

* **Sectors:** Each track is further divided into smaller, fixed-size segments called **sectors** (traditionally 512 bytes). A sector is the smallest unit of data that can be read or written.

To access a specific record, the system needs three pieces of information:

1. **Cylinder (Track) Number:** To move the arm assembly to the correct track.
2. **Surface (Head) Number:** To select which read/write head to activate.
3. **Sector Number:** To identify the correct segment on the track.

Calculating Access Time on a Hard Disk

The total time it takes to access data on a hard drive, known as **access time**, is the sum of three components:

Access Time = Seek Time + Search Time + Transfer Time

Let's break these down:

1. **Seek Time:** This is the time it takes for the arm assembly to move the read/write heads from their current position to the correct track (cylinder). **This is the slowest and most time-consuming part of the process.** Therefore, minimizing seek time is the primary goal of disk scheduling algorithms.
2. **Search Time (or Rotational Delay):** Once the head is over the correct track, this is the time it takes for the disk to rotate until the desired sector is positioned directly under the read/write head.
3. **Transfer Time:** This is the time it takes to actually transfer the data from the disk sector into the computer's main memory. This is the fastest of the three components.

Types of Hard Drives and Their Access Time Calculations

- **Fixed-Head Drives:** These are rare but have one read/write head for every single track. The heads do not move, so there is **no seek time**.
 - **Access Time = Search Time + Transfer Time**
 - The main performance factor is the rotational delay (search time). On average, the disk has to make a half rotation to find the desired sector, so the average search time is half the time of a full rotation.
- **Movable-Head Drives:** These are the standard hard drives we use, with one head per surface that moves across all tracks.
 - **Access Time = Seek Time + Search Time + Transfer Time**
 - Because seek time is the largest component, we need strategies to minimize it.

The Power of Blocking on Hard Drives

Just like with tapes, blocking records on a hard drive significantly improves performance.

- **Unblocked Records:** To read 10 individual records scattered randomly across the disk, the drive might have to perform 10 separate seeks and 10 separate searches.
 - $\text{Total Access Time} \approx 10 \times (\text{Seek Time} + \text{Search Time} + \text{Transfer Time})$
 - Using our example numbers, reading 10 unblocked records on a movable-head drive took **134.94 milliseconds**.
- **Blocked Records:** If the 10 records are stored together in one block, the drive only needs to perform **one seek** and **one search** to get to the start of the block. Then, it just transfers the entire block of data.
 - $\text{Total Access Time} \approx (1 \times \text{Seek Time}) + (1 \times \text{Search Time}) + (10 \times \text{Transfer Time for one record})$
 - Using the same example numbers, reading a block of 10 records took only **14.34 milliseconds**.

This is a massive improvement, from over 130 ms down to just 14 ms. This demonstrates why blocking is a fundamental technique for optimizing I/O performance. Of course, there's a small overhead for de-blocking the records in memory, but this is almost always negligible compared to the time saved on disk access.

The Performance Bottleneck

Over the last 40 years, the speed of CPUs and main memory has increased dramatically (think of Moore's Law). However, the mechanical nature of hard disk drives means their performance has improved at a much slower rate. This creates a significant performance gap. You can have the fastest CPU and tons of RAM, but if your system is constantly waiting for a slow hard drive, the entire system will feel sluggish. The disk subsystem is often the **bottleneck**. This is why optimizing disk access is so critical.

(At this point in the lecture, we took a short 10-minute break. Let's resume with our next topic: the strategies used to minimize that all-important seek time.)

Part 6: Device Handler Seek Strategies

Alright, welcome back. We've established that **seek time** is the biggest performance killer for movable-head hard drives. The goal of a device handler's seek strategy, or **disk scheduling algorithm**, is to minimize the total seek time by intelligently ordering the queue of I/O requests.

We will explore several of these algorithms in our next session, including First-Come, First-Served (FCFS), Shortest Seek Time First (SSTF), SCAN, and C-SCAN. For now, it's important to understand that the choice of algorithm has a direct and significant impact on the performance of the entire storage subsystem. This topic also has a strong link to our next major unit, the **File Manager**, which determines how files are physically laid out on the disk, thereby influencing the pattern of I/O requests.

Exam Preparation: Sample Question and Walkthrough

To help you prepare for the exam, let's walk through a sample question. I've given you a few hints about exam questions before, but let's look at one in detail.

Now, don't worry, the actual exam question will **never be exactly the same** as this one. But the *style* of thinking and the *type* of problem-solving will be similar. Knowing some of the questions ahead of time helps you focus, but there will always be some uncertainty. If you knew every question, you'd only study those specific answers. This way, you have a starting point but still need to explore other possibilities.

Sample Question:

In high-performance computing, large mathematical calculations can be computed more efficiently using parallel processing. Consider the following equation:

$$K = (I^J) - (4^L)^{(M/N)}$$

Your task is to break this equation into different computational steps that can be run in parallel using five CPUs. You must:

- 1. Identify the sub-tasks that can be computed concurrently.*
- 2. Use COBEGIN and COEND statements to write pseudo-code showing the parallel execution.*
- 3. Describe the intermediate values stored in temporary variables (T1, T2, etc.).*
- 4. Create a CPU allocation table assigning each sub-task to one of the five CPUs.*
- 5. Draw a parallel execution flow diagram showing how the equation is executed and how the results are merged.*

Walkthrough of a Good Answer:

A good answer will be structured step-by-step. While your specific breakdown might differ slightly, the logic should be sound. You don't have to follow my exact solution, but you must justify your steps. If you come up with a truly inspiring and novel solution that breaks from the traditional approach, that is perfectly acceptable as long as it's well-explained.

Step 1: Breakdown of the Equation

First, we identify the independent calculations that can be done in parallel.

- * $T1 = I^J$ (Exponential)
- * $T2 = 4^L$ (Exponential)
- * $T3 = M / N$ (Division)

These three calculations do not depend on each other, so they can be run on separate CPUs at the same time.

Next, we identify the dependent calculations.

- * $T4 = T2^T3$ (This depends on the results of T2 and T3)
- * $K = T1 - T4$ (The final result, which depends on T1 and T4)

Step 2: Pseudo-code with COBEGIN/COEND

This shows the parallel and sequential parts of the execution.

```
// Parallel execution of independent tasks
COBEGIN
    T1 = I ^ J;          // Executed on CPU1
    T2 = 4 ^ L;          // Executed on CPU2
    T3 = M / N;          // Executed on CPU3
COEND

// Sequential execution of dependent tasks
T4 = T2 ^ T3;           // Executed on CPU4
K = T1 - T4;            // Executed on CPU5
```


Step 3: Explanation of Intermediate Results

- * **T1:** Stores the result of I raised to the power of J .
- * **T2:** Stores the result of 4 raised to the power of L .
- * **T3:** Stores the result of M divided by N .
- * **T4:** Stores the result of T_2 raised to the power of T_3 .
- * **K:** Stores the final result of the entire equation.

Step 4: CPU Allocation Table

CPU	Task Assigned
CPU 1	Calculate $T_1 = I^J$
CPU 2	Calculate $T_2 = 4^L$
CPU 3	Calculate $T_3 = M / N$
CPU 4	Calculate $T_4 = T_2^{T_3}$ (merges T2 and T3)
CPU 5	Calculate $K = T_1 - T_4$ (merges T1 and T4)

Step 5: Parallel Execution Flow Diagram

(Here you would draw a diagram showing three parallel streams for CPU1, CPU2, and CPU3. The outputs of CPU2 and CPU3 would feed into CPU4. The output of CPU1 and CPU4 would then feed into CPU5, which produces the final result, K.)

End-of-Lecture Q&A and Final Remarks

At the end of the lecture, we had a brief Q&A session. Here are the key takeaways from that discussion:

- **On the sample question:** I've noticed from the knowledge checks that everyone in this class seems to be doing perfectly. You're all geniuses! So, I need to step up the difficulty of the exam questions to make them more challenging. (This was said with a bit of humor, but the message is: expect the exam to be challenging and require deep understanding, not just memorization.)
- **On getting more exam hints:** A student asked if I could provide more sample questions. I explained that I've already given hints on three major topics: the deadlock problem (the dining philosophers), the process scheduling calculation, and now this parallelism problem. I jokingly referred to these hints as my "Valentine's Day present" to you. Giving away too much would defeat the purpose of the exam. The goal is to give you some certainty so you don't fail, but enough uncertainty to encourage you to study the full breadth of the material.
- **On the format of the exam:** The exam will consist of multiple-choice questions and written-answer/calculation problems.
 - Some multiple-choice questions will have only **one** correct answer.
 - Others will have **multiple** correct answers (and will be marked as "select all that apply").
- **A clever question from a student:** A student asked a very insightful question that tried to anticipate the exact parameters of a calculation problem on the exam. I complimented the student on their "genius" question, as it showed they were thinking critically about how to prepare. While I couldn't confirm the exact details, it highlights the right way to think: anticipate the variables and be ready to apply the formulas and concepts to different scenarios.

If you have no further questions, that concludes our lecture. Good luck with your studies, and I hope everyone achieves a high score on the exam next Friday. See you then.

Study Tip Summary

Here is a recap of the most important points from this lecture to guide your studying for future exams and quizzes on this topic.

Key Exam Topics:

- * **The Four Functions of the Device Manager:** Monitoring, Enforcing Policies, Allocating, and Deallocating.
- * **Device Types:** Be able to clearly define and give examples of **Dedicated**, **Shared**, and **Virtual** devices. Understand the concept of **spooling**.
- * **Access Time Components:** You must know the difference between **Seek Time**, **Search Time (Rotational Delay)**, and **Transfer Time**.
- * **Access Time Calculation:** Be prepared to calculate the total access time for both **fixed-head** and **movable-head** drives.
- * **The Power of Blocking:** Understand *why* blocking is so effective for both tapes and disks. Be able to compare the access time for reading blocked vs. unblocked records.
- * **Hard Drive Anatomy:** Know the key parts: platter, track, cylinder, and sector.

Common Misunderstandings to Avoid:

- * **Seek Time vs. Search Time:** Don't confuse them. Seek time is the arm moving to the right track. Search time is the disk spinning to the right sector. Seek time is the biggest bottleneck.
- * **Old Technology is Obsolete:** Don't assume that old technologies like magnetic tape are completely gone. They are still used in niche,

large-scale applications like cloud archival storage (e.g., AWS Glacier) because they are cheap and have high capacity.

* **Forgetting Seek Time:** Remember that for movable-head drives, seek time is part of the access time calculation. For fixed-head drives, it is not.

Critical Formulas & Definitions to Memorize:

- **Access Time (Movable-Head Drive):**

$\text{Access Time} = \text{Seek Time} + \text{Search Time} + \text{Transfer Time}$

- **Access Time (Fixed-Head Drive):**

$\text{Access Time} = \text{Search Time} + \text{Transfer Time}$

- **Tape Transfer Rate:**

$\text{Transfer Rate} = \text{Tape Density (bytes/inch)} \times \text{Transport Speed (inches/second)}$

- **Key Definitions:**

- **Seek Time:** Time to move the read/write head to the correct track.
- **Search Time (Rotational Delay):** Time for the disk to rotate the correct sector under the head.
- **Cylinder:** A vertical stack of tracks at the same position on all platters.
- **Blocking:** Grouping multiple logical records into a single physical block to improve I/O efficiency.
- **Spooling:** Using a high-speed device (like a disk) as a buffer for a slow peripheral (like a printer) to free up the CPU and other resources.